

Loan Applications

1. Business Understanding

a. Introduction

In the financial industry, loan applications undergo a rigorous manual review process to determine whether they should be approved or rejected. This process involves significant labor costs and time, as human reviewers meticulously evaluate each application based on various characteristics and criteria. Given the substantial volume of applications, there is a growing need for an automated solution that can efficiently and accurately predict the outcome of loan applications, thereby reducing labor costs and improving decision-making speed.

b. Problem Statement

The manual evaluation of loan applications is costly and resource-intensive, requiring significant human labor and time. As the volume of applications grows, the need for more staff increases, escalating costs further. This process also risks variability and inconsistency in decisions. Therefore, there is a need for an automated solution to efficiently and accurately evaluate loan applications, reducing labor costs and minimizing human error.

c. Objective

Main Objective

- To create a machine learning model that would automatically reject or approve loan applications.

Specific Objectives

- Develop a machine learning model to predict the AR target variable.
- Determine optimal probability thresholds for automatically rejecting or approving loan applications to minimize labor costs and losses from incorrect decisions.
- Provide a concise description of the prediction model and the variables used.
- Summarize the business impact, including cost and loss differences, if the model were implemented in live decisions, using appropriate visualizations.

Import Libraries

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
from warnings import filterwarnings
filterwarnings(action='ignore')
```

```
# Loading the dataset
```

```
loan = pd.read_csv('Twino_data.csv')
loan.head(10)
```

	application_id	AR	payment_method	client_first_lapp_mark	\
0	1	0	METHOD1	Yes	
1	2	0	METHOD1	No	
2	3	0	METHOD1	No	
3	4	0	METHOD1	No	
4	5	0	METHOD1	No	
5	6	0	METHOD1	Yes	
6	7	1	METHOD1	Yes	
7	8	1	METHOD1	Yes	
8	9	1	METHOD1	Yes	
9	10	1	METHOD1	No	

	client_first_manual_lapp_mark	warning_count	address_check_1	\
0	Yes	2	False	
1	Yes	2	False	
2	Yes	6	False	
3	Yes	7	False	
4	No	5	False	
5	Yes	1	False	
6	Yes	11	False	
7	Yes	7	False	
8	Yes	6	False	
9	Yes	4	False	

	address_check_2	address_check_3	address_check_4	...	other_29	\
0	False	False	False	...	False	
1	False	False	False	...	False	
2	False	False	False	...	False	
3	False	False	False	...	False	
4	False	False	False	...	False	
5	False	False	False	...	False	
6	False	False	False	...	False	
7	False	False	False	...	False	
8	False	False	False	...	False	
9	False	False	False	...	False	

	data_quality_30	other_31	data_quality_32	data_quality_33	\
0	False	False	False	False	
1	False	False	False	False	
2	False	False	False	False	
3	False	False	False	False	
4	False	False	False	True	
5	False	False	False	False	

6	True	False	True	False
7	False	False	False	False
8	False	False	False	True
9	False	False	True	False

	data_quality_34	data_quality_35	other_36	accounts_check_37 \
0	False	False	False	False
1	False	False	False	False
2	False	False	False	False
3	False	False	False	False
4	False	False	False	False
5	False	False	False	False
6	False	False	False	False
7	False	False	False	False
8	False	False	False	False
9	False	False	False	False

	accounts_check_38
0	False
1	False
2	False
3	False
4	False
5	False
6	False
7	False
8	True
9	False

[10 rows x 44 columns]

a summary of the dataset
loan.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 9898 entries, 0 to 9897
Data columns (total 44 columns):
```

#	Column	Non-Null Count	Dtype
---	-----	-----	-----
0	application_id	9898 non-null	int64
1	AR	9898 non-null	int64
2	payment_method	9898 non-null	object
3	client_first_lapp_mark	9898 non-null	object
4	client_first_manual_lapp_mark	9898 non-null	object
5	warning_count	9898 non-null	int64
6	address_check_1	9898 non-null	bool
7	address_check_2	9898 non-null	bool
8	address_check_3	9898 non-null	bool
9	address_check_4	9898 non-null	bool
10	accounts_check_5	9898 non-null	bool

11	accounts_check_6	9898	non-null	bool
12	data_comparison_7	9898	non-null	bool
13	data_comparison_8	9898	non-null	bool
14	data_comparison_9	9898	non-null	bool
15	data_comparison_10	9898	non-null	bool
16	data_comparison_11	9898	non-null	bool
17	data_comparison_12	9898	non-null	bool
18	data_comparison_13	9898	non-null	bool
19	data_comparison_14	9898	non-null	bool
20	data_comparison_15	9898	non-null	bool
21	creditcard_check_16	9898	non-null	bool
22	creditcard_check_17	9898	non-null	bool
23	creditcard_check_18	9898	non-null	bool
24	data_quality_19	9898	non-null	bool
25	data_quality_20	9898	non-null	bool
26	other_21	9898	non-null	bool
27	other_22	9898	non-null	bool
28	data_quality_23	9898	non-null	bool
29	data_quality_24	9898	non-null	bool
30	data_quality_25	9898	non-null	bool
31	data_quality_26	9898	non-null	bool
32	data_quality_27	9898	non-null	bool
33	data_quality_28	9898	non-null	bool
34	other_29	9898	non-null	bool
35	data_quality_30	9898	non-null	bool
36	other_31	9898	non-null	bool
37	data_quality_32	9898	non-null	bool
38	data_quality_33	9898	non-null	bool
39	data_quality_34	9898	non-null	bool
40	data_quality_35	9898	non-null	bool
41	other_36	9898	non-null	bool
42	accounts_check_37	9898	non-null	bool
43	accounts_check_38	9898	non-null	bool

dtypes: bool(38), int64(3), object(3)
memory usage: 831.4+ KB

Data Cleaning

```
# Checking for missing values
loan.isnull().sum()
```

application_id	0
AR	0
payment_method	0
client_first_lapp_mark	0
client_first_manual_lapp_mark	0
warning_count	0
address_check_1	0
address_check_2	0

```

address_check_3      0
address_check_4      0
accounts_check_5     0
accounts_check_6     0
data_comparison_7    0
data_comparison_8    0
data_comparison_9    0
data_comparison_10   0
data_comparison_11   0
data_comparison_12   0
data_comparison_13   0
data_comparison_14   0
data_comparison_15   0
creditcard_check_16  0
creditcard_check_17  0
creditcard_check_18  0
data_quality_19      0
data_quality_20      0
other_21             0
other_22             0
data_quality_23      0
data_quality_24      0
data_quality_25      0
data_quality_26      0
data_quality_27      0
data_quality_28      0
other_29             0
data_quality_30      0
other_31             0
data_quality_32      0
data_quality_33      0
data_quality_34      0
data_quality_35      0
other_36             0
accounts_check_37    0
accounts_check_38    0
dtype: int64

```

Checking for duplicates

```

duplicates = loan.duplicated().sum()
duplicates

```

```

0

```

Dropping application_id

```

loan = loan.drop('application_id', axis=1)
loan.head()

```

AR payment_method client_first_lapp_mark

client_first_manual_lapp_mark \

0 0 METHOD1 Yes

Yes

1 0 METHOD1 No

Yes

2 0 METHOD1 No

Yes

3 0 METHOD1 No

Yes

4 0 METHOD1 No

No

warning_count address_check_1 address_check_2 address_check_3 \

0 2 False False False

1 2 False False False

2 6 False False False

3 7 False False False

4 5 False False False

address_check_4 accounts_check_5 ... other_29

data_quality_30 \

0 False False ... False False

1 False False ... False False

2 False False ... False False

3 False False ... False False

4 False False ... False False

other_31 data_quality_32 data_quality_33 data_quality_34 \

0 False False False False

1 False False False False

2 False False False False

3 False False False False

4 False False True False

data_quality_35 other_36 accounts_check_37 accounts_check_38

0 False False False False

1 False False False False

2 False False False False

3 False False False False

4 False False False False

[5 rows x 43 columns]

Exploratory Data Analysis

```
loan.describe()
```

	AR	warning_count
count	9898.000000	9898.000000
mean	0.342999	5.136593
std	0.474735	2.377723
min	0.000000	1.000000
25%	0.000000	3.000000
50%	0.000000	5.000000
75%	1.000000	7.000000
max	1.000000	16.000000

Distribution of the Target Variable (AR)

```
# Checking how many classes are there.
```

```
loan["AR"].value_counts()
```

```
AR
```

```
0    6503
```

```
1    3395
```

```
Name: count, dtype: int64
```

```
# Plotting the distribution of the target variable(AR)
```

```
plt.figure(figsize=(8, 4))
```

```
sns.countplot(x='AR', data=loan, palette='Set1')
```

```
plt.title('Distribution of Target Variable (AR)')
```

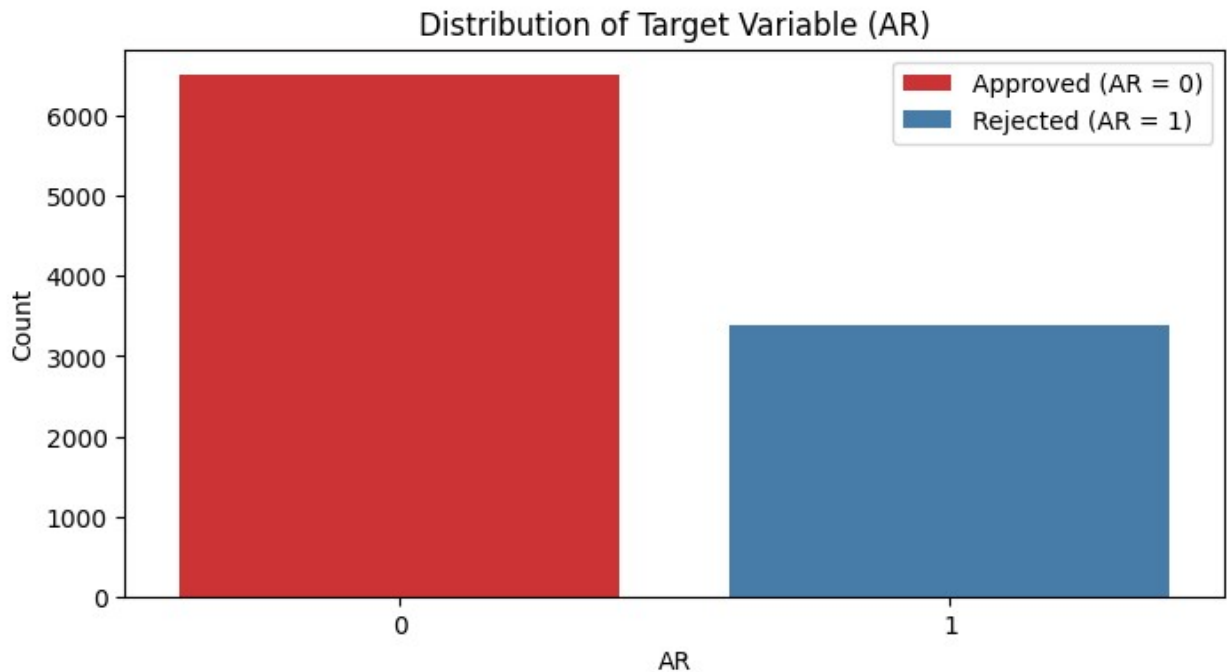
```
plt.xlabel('AR')
```

```
plt.ylabel('Count')
```

```
plt.legend(labels=['Approved (AR = 0)', 'Rejected (AR = 1)'],
```

```
loc='upper right')
```

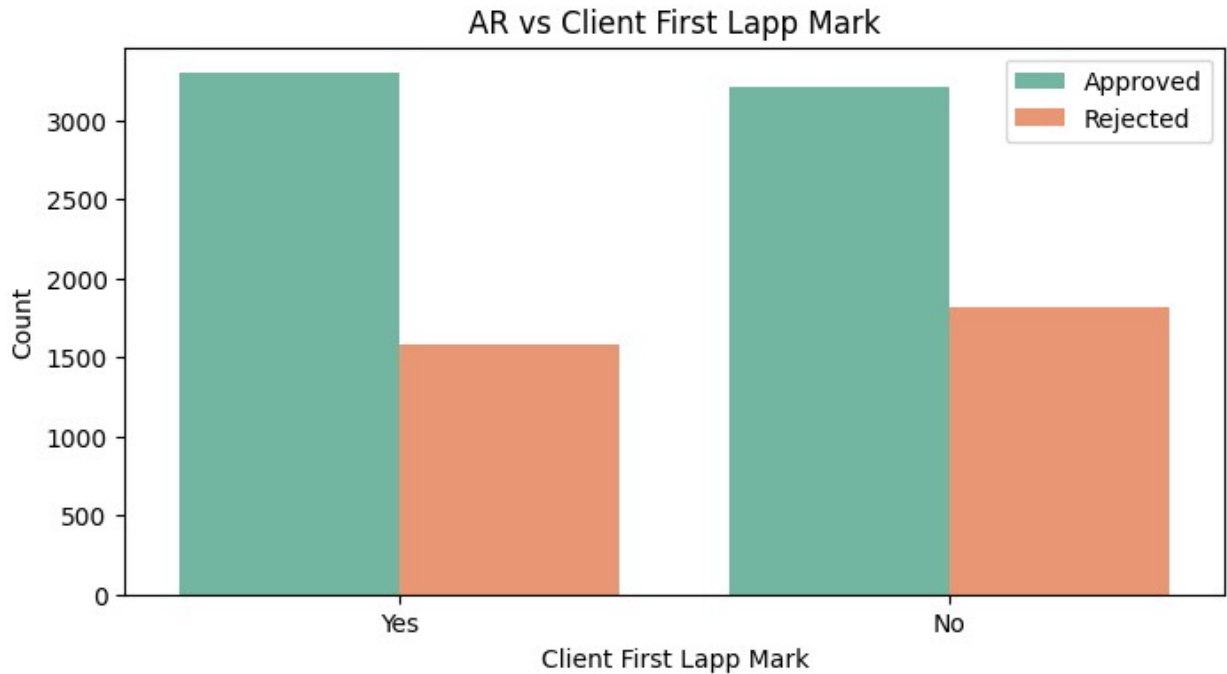
```
plt.show()
```



Bivariate Analysis

AR vs Client First Lapp Mark

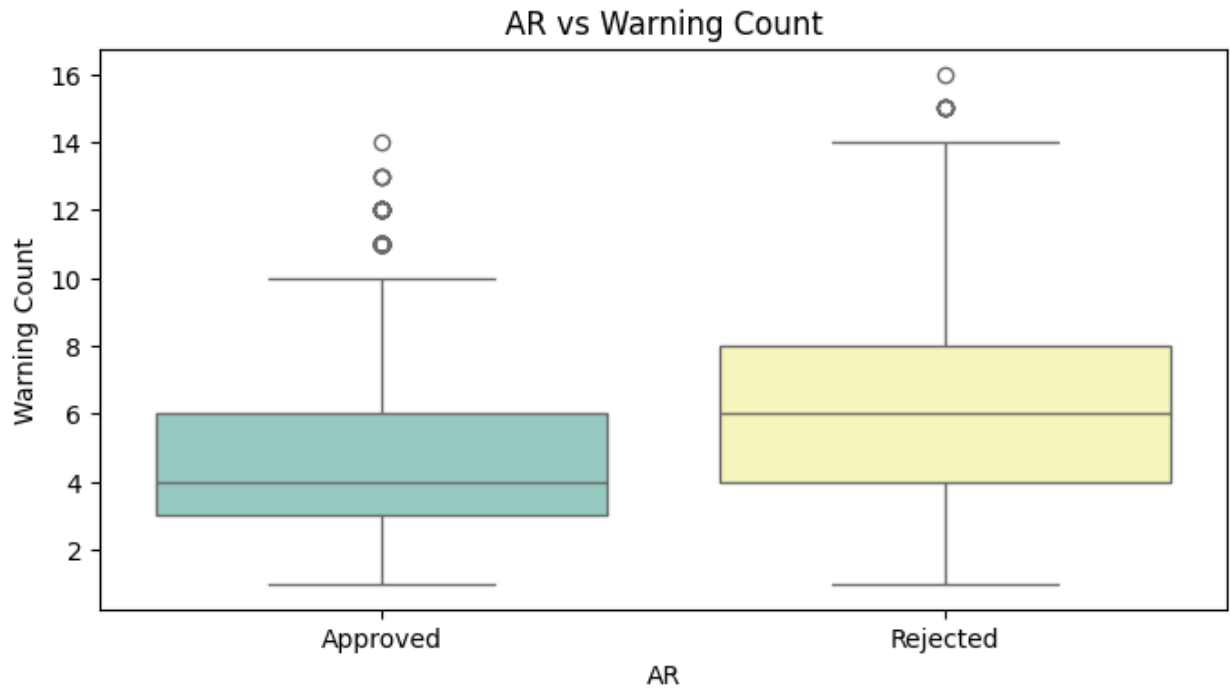
```
plt.figure(figsize=(8, 4))
sns.countplot(x='client_first_lapp_mark', hue='AR', data=loan,
palette='Set2')
plt.title('AR vs Client First Lapp Mark')
plt.xlabel('Client First Lapp Mark')
plt.ylabel('Count')
plt.legend(labels=['Approved', 'Rejected'])
plt.show()
```

- The countplot visually contrasts loan approval outcomes (AR) based on whether applicants have a prior loan application (client_first_lapp_mark), aiming to uncover actionable insights for enhancing approval processes.

AR vs Warning Count (Numerical Analysis)

```
plt.figure(figsize=(8, 4))
sns.boxplot(x='AR', y='warning_count', data=loan, palette='Set3')
plt.title('AR vs Warning Count')
plt.xlabel('AR')
plt.ylabel('Warning Count')
plt.xticks([0, 1], ['Approved', 'Rejected'])
plt.show()
```



- These box plot will give insights to optimize the loan approval process, ensuring that decisions are fair, consistent, and aligned with risk management objectives.

Data Preprocessing

Encoding

```
# columns to encode

from sklearn.preprocessing import LabelEncoder

categorical_columns = ['payment_method', 'client_first_lapp_mark',
                       'client_first_manual_lapp_mark']
boolean_columns = loan.select_dtypes(include=['bool']).columns

# coping of the DataFrame for encoding
loan_encoded = loan.copy()

# Labeling encoding for categorical columns
label_encoders = {}
for column in categorical_columns:
    le = LabelEncoder()
    loan_encoded[column] = le.fit_transform(loan_encoded[column])
    label_encoders[column] = le
```

```
# Converting boolean columns to integers
loan_encoded[boolean_columns] =
loan_encoded[boolean_columns].astype(int)
```

Handling class imbalance

```
from imblearn.over_sampling import SMOTE

# Separating features and target variable
X = loan_encoded.drop('AR', axis=1)
y = loan_encoded['AR']

# Applying SMOTE
smote = SMOTE(random_state=42)
X_resampled, y_resampled = smote.fit_resample(X, y)

# Printing the class distribution after applying SMOTE
print(y_resampled.value_counts())
```

```
AR
0    6503
1    6503
Name: count, dtype: int64
```

Train-Test Split

```
from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)
```

Feature Scaling

```
from sklearn.preprocessing import StandardScaler

scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)
```

Modelling

Baseline Model: Logistic Regression

```
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import classification_report, confusion_matrix,
accuracy_score
```

```

# Fiting the logistic regression model
log_reg = LogisticRegression(random_state=42)
log_reg.fit(X_train_scaled, y_train)

# Making predictions
y_pred = log_reg.predict(X_test_scaled)

# Calculating the confusion matrix
conf_matrix = confusion_matrix(y_test, y_pred)

# Calculating accuracy score and classification report
accuracy = accuracy_score(y_test, y_pred)
class_report = classification_report(y_test, y_pred)

# Printing accuracy score and classification report
print("Accuracy Score:", accuracy)
print("\nClassification Report:\n", class_report)

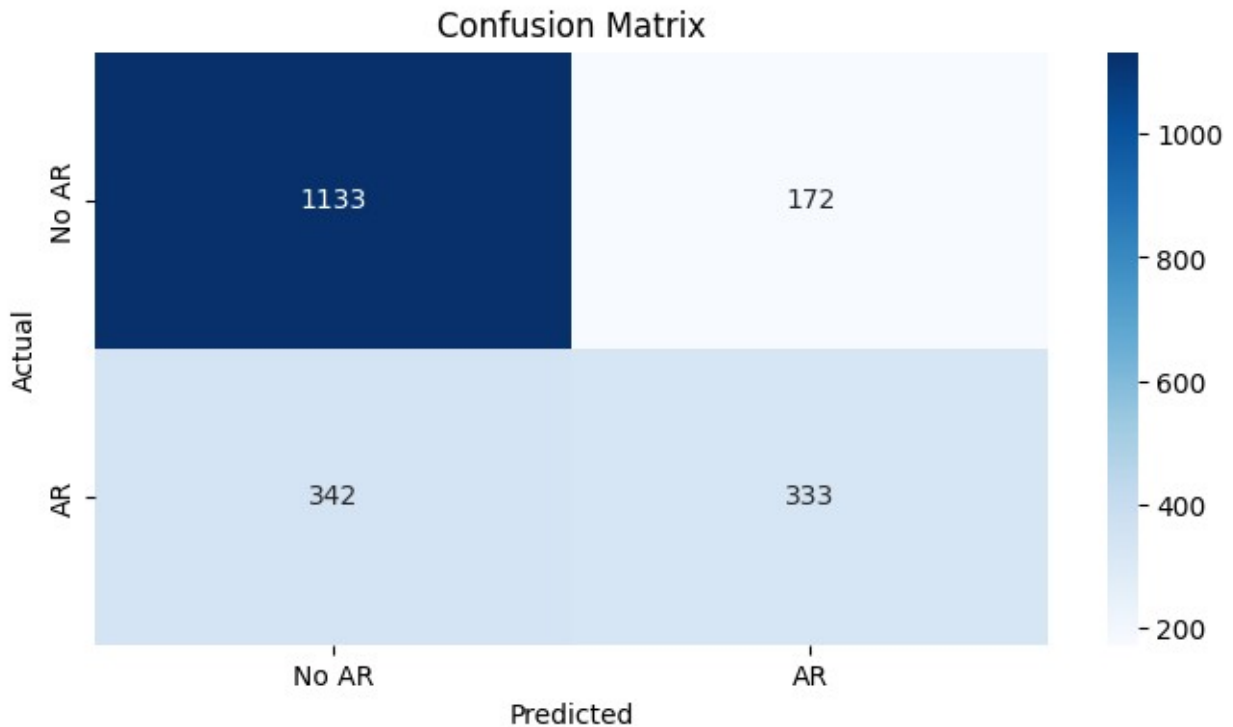
# Ploting the confusion matrix
plt.figure(figsize=(8, 4))
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues',
            xticklabels=['No AR', 'AR'], yticklabels=['No AR', 'AR'])
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Confusion Matrix')
plt.show()

```

Accuracy Score: 0.7404040404040404

Classification Report:

	precision	recall	f1-score	support
0	0.77	0.87	0.82	1305
1	0.66	0.49	0.56	675
accuracy			0.74	1980
macro avg	0.71	0.68	0.69	1980
weighted avg	0.73	0.74	0.73	1980



```
# Training accuracy
train_accuracy = accuracy_score(y_train,
log_reg.predict(X_train_scaled))
print(f"Training Accuracy: {train_accuracy:.4f}")

# Test accuracy
test_accuracy = accuracy_score(y_test, y_pred)
print(f"Test Accuracy: {test_accuracy:.4f}")

Training Accuracy: 0.7267
Test Accuracy: 0.7404
```

Random Forest

```
from sklearn.ensemble import RandomForestClassifier

# Initializing the Random Forest classifier
rf_clf = RandomForestClassifier(random_state=42)

# Fitting the model on the training data
rf_clf.fit(X_train, y_train)

# Predicting on the test data
y_pred = rf_clf.predict(X_test)

# Evaluating the model
```

```

accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy Score: {accuracy:.4f}\n")

conf_matrix = confusion_matrix(y_test, y_pred)
print(f"Confusion Matrix:\n {conf_matrix}\n")

class_report = classification_report(y_test, y_pred)
print(f"Classification Report:\n {class_report}")

```

Accuracy Score: 0.7217

Confusion Matrix:

```

[[1078  227]
 [ 324  351]]

```

Classification Report:

	precision	recall	f1-score	support
0	0.77	0.83	0.80	1305
1	0.61	0.52	0.56	675
accuracy			0.72	1980
macro avg	0.69	0.67	0.68	1980
weighted avg	0.71	0.72	0.72	1980

Support Vector Machine

```

from sklearn.svm import SVC

# Initializing the SVM classifier
svm_clf = SVC(kernel='rbf', random_state=42)

# Fiting the model on the training data
svm_clf.fit(X_train, y_train)

# Predicting on the test data
y_pred = svm_clf.predict(X_test)

# Evaluating the model
accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy Score: {accuracy:.4f}\n")

conf_matrix = confusion_matrix(y_test, y_pred)
print(f"Confusion Matrix:\n {conf_matrix}\n")

class_report = classification_report(y_test, y_pred)
print(f"Classification Report:\n {class_report}")

```

Accuracy Score: 0.7354

Confusion Matrix:

```
[[1125  180]
 [ 344  331]]
```

Classification Report:

	precision	recall	f1-score	support
0	0.77	0.86	0.81	1305
1	0.65	0.49	0.56	675
accuracy			0.74	1980
macro avg	0.71	0.68	0.68	1980
weighted avg	0.73	0.74	0.72	1980

XG BOOST

```
import xgboost as xgb
```

```
# Initializing the XGBClassifier
```

```
xgb_clf = xgb.XGBClassifier(random_state=42)
```

```
# Fitting the model on the training data
```

```
xgb_clf.fit(X_train, y_train)
```

```
# Predicting on the test data
```

```
y_pred = xgb_clf.predict(X_test)
```

```
# Evaluating the model
```

```
accuracy = accuracy_score(y_test, y_pred)
```

```
print(f"Accuracy Score: {accuracy:.4f}\n")
```

```
conf_matrix = confusion_matrix(y_test, y_pred)
```

```
print(f"Confusion Matrix:\n {conf_matrix}\n")
```

```
class_report = classification_report(y_test, y_pred)
```

```
print(f"Classification Report:\n {class_report}")
```

Accuracy Score: 0.7389

Confusion Matrix:

```
[[1111  194]
 [ 323  352]]
```

Classification Report:

	precision	recall	f1-score	support
0	0.77	0.85	0.81	1305

	1	0.64	0.52	0.58	675
accuracy				0.74	1980
macro avg		0.71	0.69	0.69	1980
weighted avg		0.73	0.74	0.73	1980

Hyperparameter Tuning (Logistic Regression)

```

from sklearn.model_selection import GridSearchCV

# Defining the parameters grid
param_grid = {
    'C': [0.001, 0.01, 0.1, 1, 10, 100], # Regularization parameter
    'penalty': ['l1', 'l2'], # Regularization penalty
    ('l1' or 'l2')
    'solver': ['liblinear', 'saga'] # Optimization algorithm
}

# Initializing GridSearchCV
grid_search =
GridSearchCV(estimator=LogisticRegression(random_state=42),
              param_grid=param_grid,
              cv=5, # 5-fold cross-validation
              scoring='accuracy',
              verbose=1,
              n_jobs=-1)

# Fitting GridSearchCV
grid_search.fit(X_train_scaled, y_train)

# Best parameters and best score
print("Best Parameters:", grid_search.best_params_)
print("Best CV Score:", grid_search.best_score_)

Fitting 5 folds for each of 24 candidates, totalling 120 fits
Best Parameters: {'C': 0.1, 'penalty': 'l1', 'solver': 'liblinear'}
Best CV Score: 0.7230372263379212

# Using the best parameters found by GridSearchCV
best_log_reg = grid_search.best_estimator_

# Retraining with best parameters
best_log_reg.fit(X_train_scaled, y_train)

# Predicting on test data
y_pred_tuned = best_log_reg.predict(X_test_scaled)

# Evaluating the tuned model
print("Tuned Model Accuracy Score:", accuracy_score(y_test,
y_pred_tuned))

```



```
print("\nTuned Model Confusion Matrix:\n", confusion_matrix(y_test,
y_pred_tuned))
print("\nTuned Model Classification Report:\n",
classification_report(y_test, y_pred_tuned))
```

Tuned Model Accuracy Score: 0.7404040404040404

Tuned Model Confusion Matrix:

```
[[1136  169]
 [ 345  330]]
```

Tuned Model Classification Report:

	precision	recall	f1-score	support
0	0.77	0.87	0.82	1305
1	0.66	0.49	0.56	675
accuracy			0.74	1980
macro avg	0.71	0.68	0.69	1980
weighted avg	0.73	0.74	0.73	1980

The best performing model is Logistic Regression Model

Calculating the threshold

```
# Get predicted probabilities
y_prob = log_reg.predict_proba(X_test_scaled)[: , 1]

# Initializing the variables
thresholds = np.arange(0.0, 1.0, 0.01)
best_threshold = 0
min_cost = float('inf')

for threshold in thresholds:
    # Getting predicted labels based on threshold
    y_pred_threshold = (y_prob >= threshold).astype(int)

    # Calculating confusion matrix
    TN, FP, FN, TP = confusion_matrix(y_test,
y_pred_threshold).ravel()

    # Calculating costs
    labor_cost = 5 * len(y_test) # Cost for processing all
applications
    loss_FP = 17 * FP
    loss_FN = 14 * FN
```

```

total_cost = labor_cost + loss_FP + loss_FN

# Checking if this is the best threshold
if total_cost < min_cost:
    min_cost = total_cost
    best_threshold = threshold

print(f'Best Threshold: {best_threshold}, Minimum Cost: {min_cost}
EUR')

```

Best Threshold: 0.59, Minimum Cost: 17448 EUR

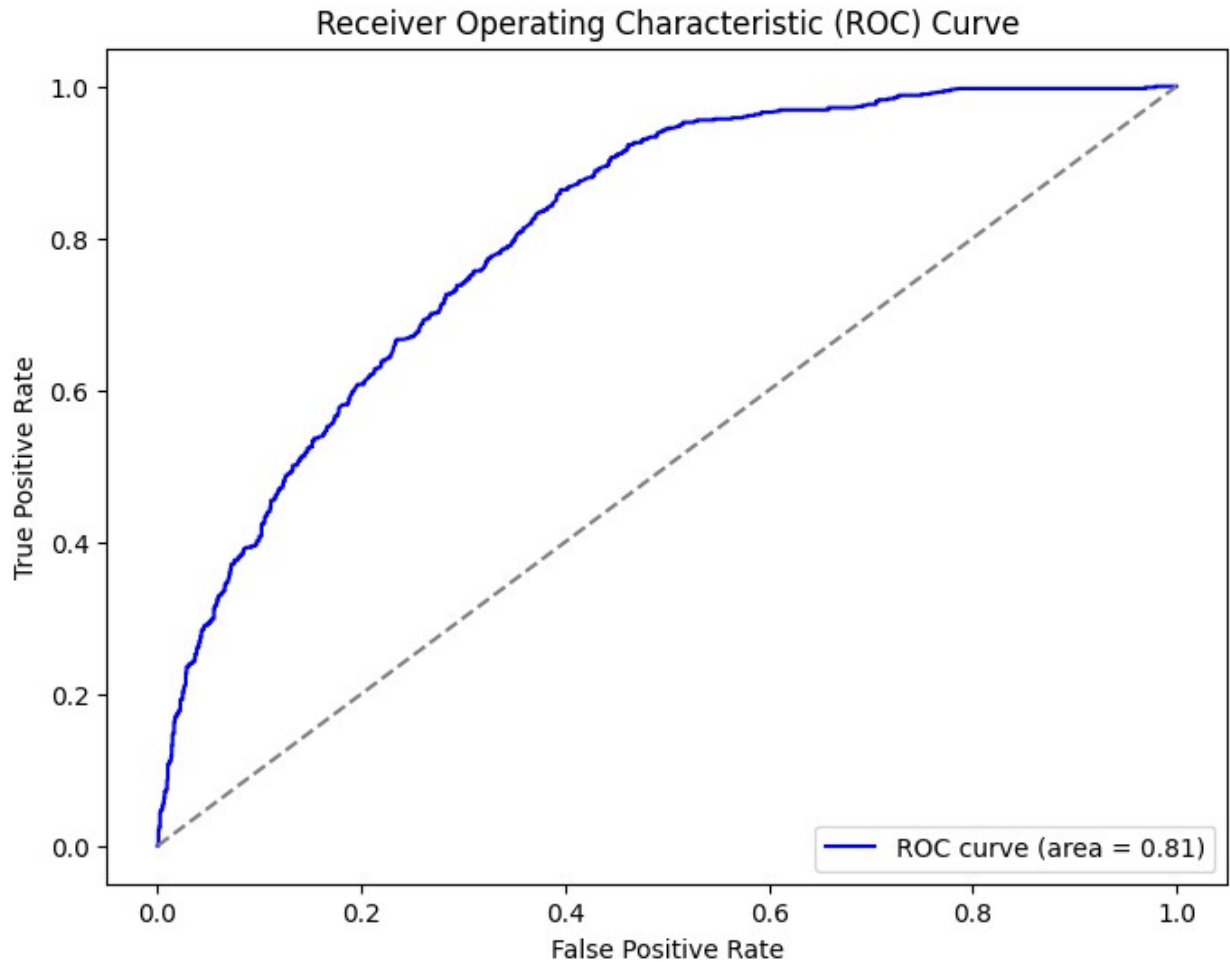
```

from sklearn.metrics import roc_curve, roc_auc_score

# ROC Curve
fpr, tpr, thresholds = roc_curve(y_test, y_prob)
roc_auc = roc_auc_score(y_test, y_prob)

plt.figure(figsize=(8, 6))
plt.plot(fpr, tpr, color='blue', label=f'ROC curve (area =
{roc_auc:.2f})')
plt.plot([0, 1], [0, 1], color='grey', linestyle='--')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve')
plt.legend(loc='lower right')
plt.show()

```



- An AUC of 0.81 suggests that there is a 81% chance that the model will be able to distinguish between positive and negative classes.

Cost vs Threshold

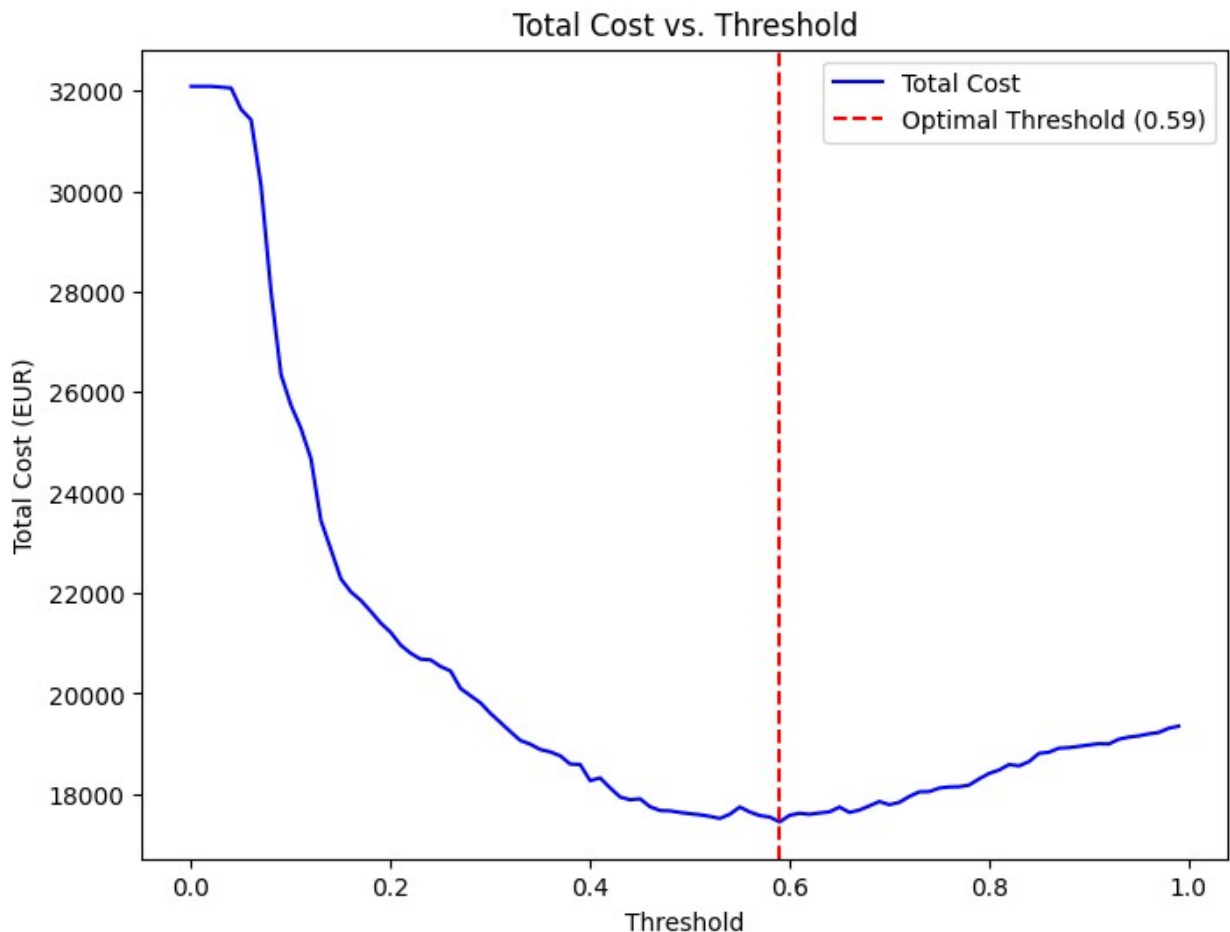
```
# Cost vs. Threshold Plot
costs = []
for threshold in np.arange(0.0, 1.0, 0.01):
    y_pred_threshold = (y_prob >= threshold).astype(int)
    TN, FP, FN, TP = confusion_matrix(y_test,
    y_pred_threshold).ravel()
    labor_cost = 5 * len(y_test)
    loss_FP = 17 * FP
    loss_FN = 14 * FN
    total_cost = labor_cost + loss_FP + loss_FN
    costs.append(total_cost)

plt.figure(figsize=(8, 6))
plt.plot(np.arange(0.0, 1.0, 0.01), costs, color='blue', label='Total
Cost')
plt.axvline(x=0.59, color='red', linestyle='--', label='Optimal
```

```

Threshold (0.59)')
plt.xlabel('Threshold')
plt.ylabel('Total Cost (EUR)')
plt.title('Total Cost vs. Threshold')
plt.legend(loc='upper right')
plt.show()

```



Business Impact Analysis

```

# Costs with the model (using optimal threshold)
labor_cost_model = 5 * (TP + FP + TN + FN) # Total labor cost for all
applications

# costs per instance
cost_FP = 17 # Cost per false positive
cost_FN = 14 # Cost per false negative

total_cost_model = labor_cost_model + (cost_FP * FP) + (cost_FN * FN)

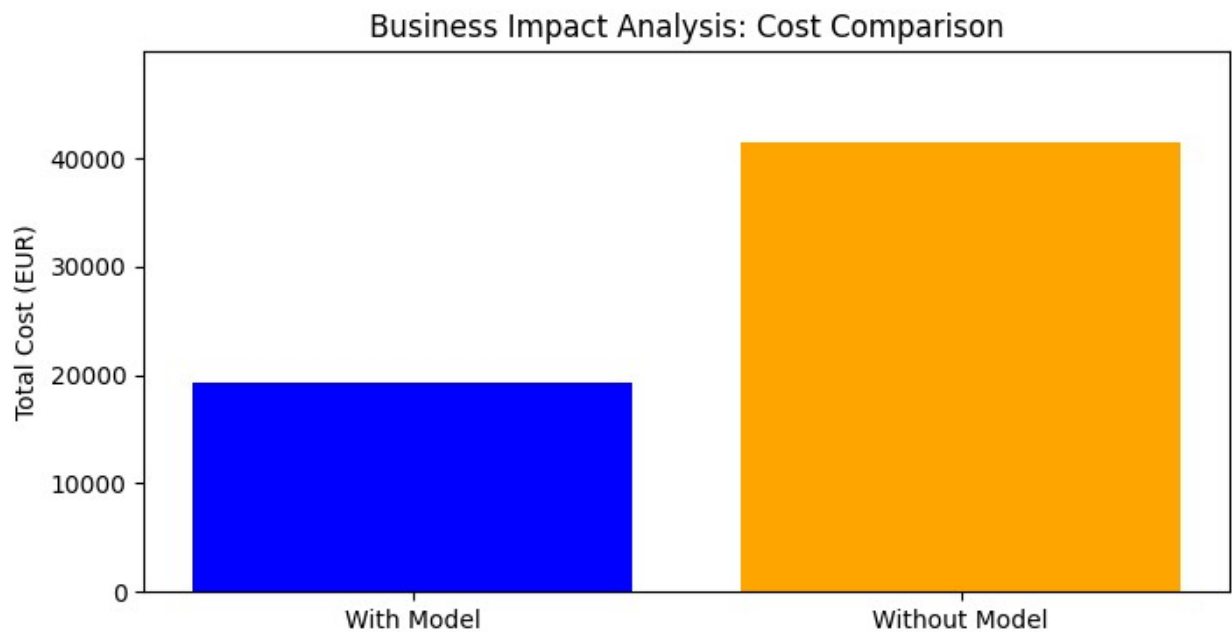
# Costs without the model (processing all applications)
total_cost_without_model = labor_cost_model + (cost_FP * (FP + TN)) +

```

```
(cost_FN * (FN + TP))

# Visualizing the comparison
categories = ['With Model', 'Without Model']
costs = [total_cost_model, total_cost_without_model]

plt.figure(figsize=(8, 4))
plt.bar(categories, costs, color=['blue', 'orange'])
plt.ylabel('Total Cost (EUR)')
plt.title('Business Impact Analysis: Cost Comparison')
plt.ylim(0, max(costs) * 1.2)
plt.show()
```



The visualization vividly demonstrates the financial advantages of deploying the Logistic Regression Model for loan approval decisions